

```
In [49]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import confusion_matrix, f1_score
from sklearn.metrics import classification_report

# Load the dataset
url = "https://raw.githubusercontent.com/tejavkat473/Machine-Learning/main/winequality-white.csv"
data = pd.read_csv(url)
data
```

Out[49]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
<b>0</b>	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.00100	3.00	0.45	8.8	6
<b>1</b>	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.99400	3.30	0.49	9.5	6
<b>2</b>	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.99510	3.26	0.44	10.1	6
<b>3</b>	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.40	9.9	6
<b>4</b>	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.99560	3.19	0.40	9.9	6
...	...	...	...	...	...	...	...	...	...	...	...	...
<b>4893</b>	6.2	0.21	0.29	1.6	0.039	24.0	92.0	0.99114	3.27	0.50	11.2	6
<b>4894</b>	6.6	0.32	0.36	8.0	0.047	57.0	168.0	0.99490	3.15	0.46	9.6	5
<b>4895</b>	6.5	0.24	0.19	1.2	0.041	30.0	111.0	0.99254	2.99	0.46	9.4	6
<b>4896</b>	5.5	0.29	0.30	1.1	0.022	20.0	110.0	0.98869	3.34	0.38	12.8	7
<b>4897</b>	6.0	0.21	0.38	0.8	0.020	22.0	98.0	0.98941	3.26	0.32	11.8	6

4898 rows × 12 columns

```
In [50]: print(data.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4898 entries, 0 to 4897
Data columns (total 12 columns):
 #   Column              Non-Null Count  Dtype  
---  -
 0   fixed acidity       4898 non-null   float64
 1   volatile acidity    4898 non-null   float64
 2   citric acid         4898 non-null   float64
 3   residual sugar      4898 non-null   float64
 4   chlorides           4898 non-null   float64
 5   free sulfur dioxide 4898 non-null   float64
 6   total sulfur dioxide 4898 non-null   float64
 7   density             4898 non-null   float64
 8   pH                 4898 non-null   float64
 9   sulphates           4898 non-null   float64
10   alcohol             4898 non-null   float64
11   quality             4898 non-null   int64  
dtypes: float64(11), int64(1)
memory usage: 459.3 KB
None
```

```
In [51]: print(data.isnull().sum())
```

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                0
sulphates          0
alcohol            0
quality            0
dtype: int64
```

```
In [52]: # Total number of missing values
print("Total number of missing values:", data.isnull().sum().sum())
```

```
Total number of missing values: 0
```

```
In [53]: print(data.shape)
```

(4898, 12)

```
In [54]: duplicate = data.duplicated()
print(duplicate.sum())
```

937

```
In [55]: data.drop_duplicates(inplace=True)
```

```
In [56]: duplicate = data.duplicated()
print(duplicate.sum())
```

0

```
In [57]: print(data.shape)
```

(3961, 12)

```
In [58]: data.describe()
```

Out[58]:


	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates
<b>count</b>	3961.000000	3961.000000	3961.000000	3961.000000	3961.000000	3961.000000	3961.000000	3961.000000	3961.000000	3961.000000
<b>mean</b>	6.839346	0.280538	0.334332	5.914819	0.045905	34.889169	137.193512	0.993790	3.195458	0.490351
<b>std</b>	0.866860	0.103437	0.122446	4.861646	0.023103	17.210021	43.129065	0.002905	0.151546	0.113523
<b>min</b>	3.800000	0.080000	0.000000	0.600000	0.009000	2.000000	9.000000	0.987110	2.720000	0.220000
<b>25%</b>	6.300000	0.210000	0.270000	1.600000	0.035000	23.000000	106.000000	0.991620	3.090000	0.410000
<b>50%</b>	6.800000	0.260000	0.320000	4.700000	0.042000	33.000000	133.000000	0.993500	3.180000	0.480000
<b>75%</b>	7.300000	0.330000	0.390000	8.900000	0.050000	45.000000	166.000000	0.995710	3.290000	0.550000
<b>max</b>	14.200000	1.100000	1.660000	65.800000	0.346000	289.000000	440.000000	1.038980	3.820000	1.080000

```
In [59]: data.corr()
```

Out[59]:


	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol	quality
<b>fixed acidity</b>	1.000000	-0.019214	0.298959	0.083620	0.024036	-0.058396	0.082425	0.266091	-0.431274	-0.017453	-0.110788	-0.124636
<b>volatile acidity</b>	-0.019214	1.000000	-0.163228	0.098340	0.086287	-0.102471	0.102315	0.060603	-0.046954	-0.021150	0.046815	-0.190678
<b>citric acid</b>	0.298959	-0.163228	1.000000	0.106269	0.132590	0.091681	0.122845	0.160076	-0.183015	0.049442	-0.076514	0.007065
<b>residual sugar</b>	0.083620	0.098340	0.106269	1.000000	0.076091	0.306835	0.409583	0.820498	-0.165997	-0.020503	-0.398167	-0.117339
<b>chlorides</b>	0.024036	0.086287	0.132590	0.076091	1.000000	0.101272	0.191145	0.253088	-0.090573	0.017871	-0.356928	-0.217739
<b>free sulfur dioxide</b>	-0.058396	-0.102471	0.091681	0.306835	0.101272	1.000000	0.619437	0.294638	-0.007750	0.037932	-0.251768	0.010507
<b>total sulfur dioxide</b>	0.082425	0.102315	0.122845	0.409583	0.191145	0.619437	1.000000	0.536868	0.008239	0.136544	-0.446643	-0.183356
<b>density</b>	0.266091	0.060603	0.160076	0.820498	0.253088	0.294638	0.536868	1.000000	-0.063734	0.082048	-0.760162	-0.337805
<b>pH</b>	-0.431274	-0.046954	-0.183015	-0.165997	-0.090573	-0.007750	0.008239	-0.063734	1.000000	0.142353	0.093095	0.123829
<b>sulphates</b>	-0.017453	-0.021150	0.049442	-0.020503	0.017871	0.037932	0.136544	0.082048	0.142353	1.000000	-0.022850	0.053200
<b>alcohol</b>	-0.110788	0.046815	-0.076514	-0.398167	-0.356928	-0.251768	-0.446643	-0.760162	0.093095	-0.022850	1.000000	0.462869
<b>quality</b>	-0.124636	-0.190678	0.007065	-0.117339	-0.217739	0.010507	-0.183356	-0.337805	0.123829	0.053200	0.462869	1.000000

In [60]: `sns.heatmap(data.corr(), center=0)`Out[60]: `<AxesSubplot:>`



```
In [61]: # Calculate the correlation matrix
corr_matrix = data.corr()

# Display the correlation matrix
plt.figure(figsize=(12, 8))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', linewidths=0.5)
plt.title('Correlation Matrix')
plt.show()
```



```
In [95]: import pandas as pd
import matplotlib.pyplot as plt
```


```
# Import the dataset
url = "https://raw.githubusercontent.com/tejavenkat473/Machine-Learning/main/winequality-white.csv"
data = pd.read_csv(url)

# Specify the target variable
target_variable = 'quality' # Assuming 'quality' is the column to predict

# Count the occurrences of each class
class_counts = data[target_variable].value_counts()

# Plot the class distribution
plt.figure(figsize=(4, 3))
plt.bar(class_counts.index, class_counts.values)
plt.xlabel('Class')
plt.ylabel('Count')
plt.title('Class Distribution')
plt.show()

# Check if class imbalance exists
is_imbalanced = any(class_counts.values != class_counts.values[0])
if is_imbalanced:
    print("Class imbalance exists.")
else:
    print("Class imbalance does not exist.")
```



Class imbalance exists.

```
In [96]: import pandas as pd
import matplotlib.pyplot as plt
from imblearn.under_sampling import RandomUnderSampler
from imblearn.over_sampling import RandomOverSampler, SMOTE

# Import the dataset
url = "https://raw.githubusercontent.com/tejavenkat473/Machine-Learning/main/winequality-white.csv"
data = pd.read_csv(url)

# Specify the target variable
target_variable = 'quality' # Assuming 'quality' is the column to predict

# Count the occurrences of each class before sampling
class_counts_before = data[target_variable].value_counts()

# Plot the class distribution before sampling
plt.figure(figsize=(4, 3))
plt.bar(class_counts_before.index, class_counts_before.values)
plt.xlabel('Class')
plt.ylabel('Count')
plt.title('Class Distribution (Before Sampling)')
plt.show()

# Perform Random Under-sampling
rus = RandomUnderSampler(random_state=42)
X_resampled_under, y_resampled_under = rus.fit_resample(data.drop(target_variable, axis=1), data[target_variable])

# Count the occurrences of each class after Random Under-sampling
class_counts_under = pd.Series(y_resampled_under).value_counts()


# Plot the class distribution after Random Under-sampling
plt.figure(figsize=(4, 3))
plt.bar(class_counts_under.index, class_counts_under.values)
plt.xlabel('Class')
plt.ylabel('Count')
plt.title('Class Distribution (After Random Under-sampling)')
plt.show()

# Perform Random Over-sampling
ros = RandomOverSampler(random_state=42)
X_resampled_over, y_resampled_over = ros.fit_resample(data.drop(target_variable, axis=1), data[target_variable])


# Count the occurrences of each class after Random Over-sampling
class_counts_over = pd.Series(y_resampled_over).value_counts()
```




```
# Plot the class distribution after Random Over-sampling
plt.figure(figsize=(4, 3))
plt.bar(class_counts_over.index, class_counts_over.values)
plt.xlabel('Class')
plt.ylabel('Count')
plt.title('Class Distribution (After SMOTE Over-sampling)')
plt.show()
```



Class Distribution (After Random Under-sampling)



Class Distribution (After SMOTE Over-sampling)



```
In [97]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score
```

```

url = "https://raw.githubusercontent.com/tejavenkat473/Machine-Learning/main/winequality-white.csv"
data = pd.read_csv(url)

# Preprocessing
X = data.drop('quality', axis=1) # Assuming 'target_variable' is the column to predict
y = data['quality']

# Convert categorical variables to numerical using Label encoding
label_encoder = LabelEncoder()
X_encoded = X.apply(label_encoder.fit_transform)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)

dt_model = DecisionTreeClassifier()
dt_model.fit(X_train_scaled, y_train)
dt_predictions = dt_model.predict(X_test_scaled)
dt_accuracy = accuracy_score(y_test, dt_predictions)

print("Decision Tree Accuracy:", dt_accuracy)

```

Decision Tree Accuracy: 0.6051020408163266

```

In [98]: # Calculate confusion matrix
dt_cm = confusion_matrix(y_test, dt_predictions)

# Calculate F1 score
dt_f1 = f1_score(y_test, dt_predictions, average='weighted')

print("Decision Tree Confusion Matrix:")
print(dt_cm)
print("Decision Tree F1 Score:", dt_f1)

```

Decision Tree Confusion Matrix:

```

[[ 0  1  2  2  0  0  0]
 [ 1  7  7  7  2  1  0]
 [ 0 14 185  81  9  2  0]
 [ 1 10  80 277  51 12  1]
 [ 3  2  5  56 108 17  1]
 [ 0  1  0  7 11 16  0]
 [ 0  0  0  0  0  0  0]]

```

Decision Tree F1 Score: 0.6092344123362659

```
In [99]: # Generate classification report
dt_classification_report = classification_report(y_test, dt_predictions)

print("Decision Tree Classification Report:")
print(dt_classification_report)
```

```
Decision Tree Classification Report:
              precision    recall  f1-score   support

     3         0.00         0.00         0.00         5
     4         0.20         0.28         0.23        25
     5         0.66         0.64         0.65       291
     6         0.64         0.64         0.64       432
     7         0.60         0.56         0.58       192
     8         0.33         0.46         0.39        35
     9         0.00         0.00         0.00         0

 accuracy          0.61       980
 macro avg         0.35         0.37         0.36       980
 weighted avg         0.61         0.61         0.61       980
```

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Recall and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
In [7]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

url = "https://raw.githubusercontent.com/tejavenkat473/Machine-Learning/main/winequality-white.csv"
data = pd.read_csv(url)
```

```
# Preprocessing
X = data.drop('quality', axis=1) # Assuming 'target_variable' is the column to predict
y = data['quality']

# Convert categorical variables to numerical using Label encoding
label_encoder = LabelEncoder()
X_encoded = X.apply(label_encoder.fit_transform)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)

# Create instances of the classification algorithms
lr_model = LogisticRegression()
lr_model.fit(X_train, y_train)
lr_predictions = lr_model.predict(X_test)
lr_accuracy = accuracy_score(y_test, lr_predictions)
print("Logistic Regression Accuracy:", lr_accuracy)
```

Logistic Regression Accuracy: 0.523469387755102

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

In [104...

```
# Calculate confusion matrix
lr_cm = confusion_matrix(y_test, lr_predictions)

# Calculate F1 score
lr_f1 = f1_score(y_test, lr_predictions, average='weighted')

print("Logistic Regression Confusion Matrix:")
print(lr_cm)
print("Logistic Regression F1 Score:", lr_f1)
```

Logistic Regression Confusion Matrix:

```
[[ 0  0  2  2  0  1  0]
 [ 0  1 12 12  0  0  0]
 [ 0  1 151 137  1  1  0]
 [ 0  0  80 322 29  0  1]
 [ 0  0  13 133 46  0  0]
 [ 0  0  2  25  8  0  0]
 [ 0  0  0  0  0  0  0]]
```

Logistic Regression F1 Score: 0.4970071183496042

In [102...

```
# Generate classification report
lr_classification_report = classification_report(y_test, lr_predictions)

print("Logistic Regression Classification Report:")
print(lr_classification_report)
```

Logistic Regression Classification Report:

	precision	recall	f1-score	support
3	0.00	0.00	0.00	5
4	0.50	0.04	0.07	25
5	0.58	0.52	0.55	291
6	0.51	0.75	0.61	432
7	0.55	0.24	0.33	192
8	0.00	0.00	0.00	35
9	0.00	0.00	0.00	0
accuracy			0.53	980
macro avg	0.31	0.22	0.22	980
weighted avg	0.52	0.53	0.50	980

```

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Pre
cision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parame
ter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Rec
all and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to c
ontrol this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Pre
cision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parame
ter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Rec
all and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to c
ontrol this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Pre
cision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parame
ter to control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\_classification.py:1318: UndefinedMetricWarning: Rec
all and F-score are ill-defined and being set to 0.0 in labels with no true samples. Use `zero_division` parameter to c
ontrol this behavior.
    _warn_prf(average, modifier, msg_start, len(result))

```

In [119...

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.preprocessing import StandardScaler

# Load the dataset
url = "https://raw.githubusercontent.com/tejavenkat473/Machine-Learning/main/winequality-white.csv"
df = pd.read_csv(url)

# Separate the features (X) and target variable (y)
X = df.drop('quality', axis=1)
y = df['quality']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize the features
scaler = StandardScaler()

```

```
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)

def plot_decision_boundary(model, X, y):
    # Create a meshgrid of feature values
    h = 0.02 # step size in the mesh
    x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, h),
                          np.arange(y_min, y_max, h))

    # Make predictions on the meshgrid points
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    # Plot the decision boundaries and data points
    plt.contourf(xx, yy, Z, alpha=0.8)
    plt.scatter(X[:, 0], X[:, 1], c=y, cmap=plt.cm.Paired)
    plt.xlabel('Feature 1')
    plt.ylabel('Feature 2')
    plt.title('Decision Boundaries')
    plt.show()

    # Create an SVM classifier with a linear kernel
    svm_linear = SVC(kernel='linear')
    svm_linear.fit(X_train[:, :2], y_train)
    plot_decision_boundary(svm_linear, X_train[:, :2], y_train)

    # Create an SVM classifier with a polynomial kernel
    svm_poly = SVC(kernel='poly', degree=3)
    svm_poly.fit(X_train[:, :2], y_train)
    plot_decision_boundary(svm_poly, X_train[:, :2], y_train)

    # Create an SVM classifier with an RBF kernel
    svm_rbf = SVC(kernel='rbf')
    svm_rbf.fit(X_train[:, :2], y_train)
    plot_decision_boundary(svm_rbf, X_train[:, :2], y_train)

    # Standardize the features
    scaler = StandardScaler()
    X_train = scaler.fit_transform(X_train)
    X_test = scaler.transform(X_test)


    # Create an SVM classifier with a kernel method
    svm_model = SVC(kernel='rbf')
```




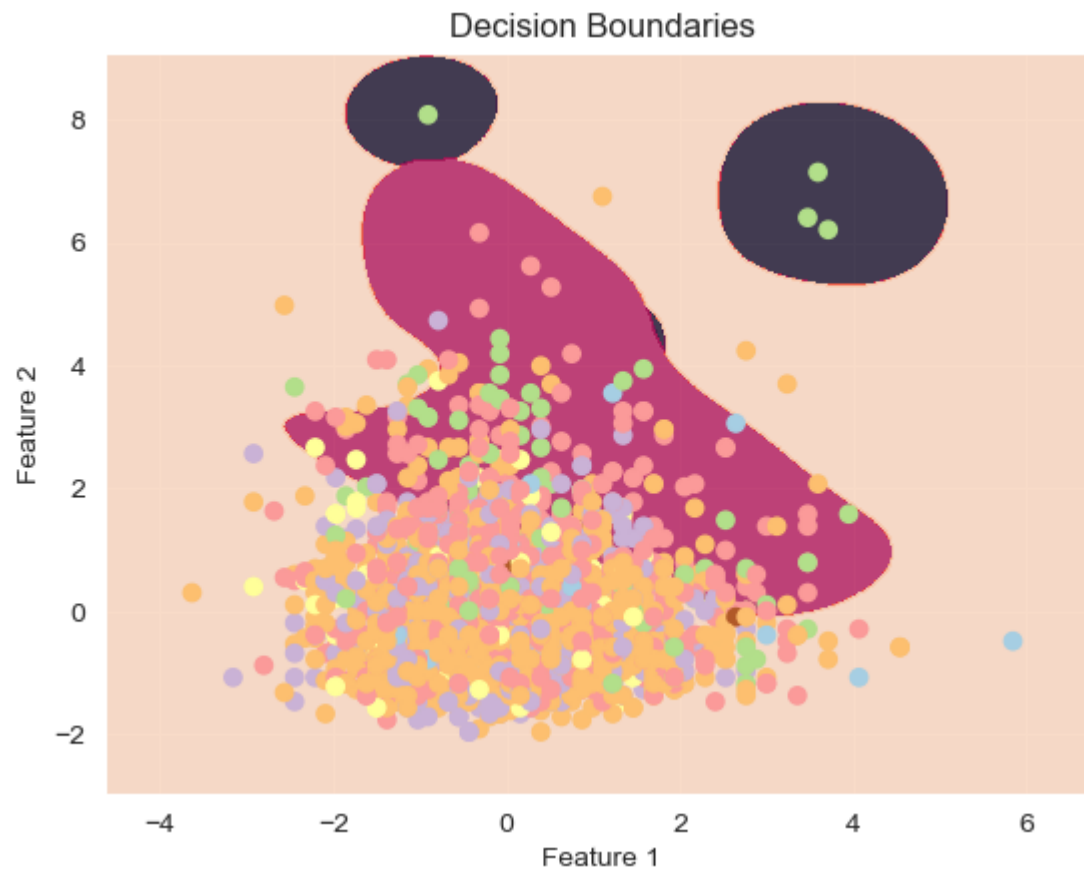
```
# Train the SVM model
svm_model.fit(X_train, y_train)

# Make predictions on the test set
y_pred = svm_model.predict(X_test)

# Evaluate the model's accuracy
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
```







Accuracy: 0.5612244897959183

In [122...

```
# Calculate confusion matrix
svm_cm = confusion_matrix(y_test, svm_predictions)

# Calculate F1 score
svm_f1 = f1_score(y_test, svm_predictions, average='weighted')

print("SVM Confusion Matrix:")
print(svm_cm)
print("SVM F1 Score:", svm_f1)
```

```
SVM Confusion Matrix:
[[ 0  0  1  4  0  0]
 [ 0  2 15  8  0  0]
 [ 0  3 167 120  1  0]
 [ 0  0  82 333 17  0]
 [ 0  0  5 139 48  0]
 [ 0  0  0  27  8  0]]
SVM F1 Score: 0.5270798963496123
```

In [121]...

```
# Generate classification report
svm_classification_report = classification_report(y_test, svm_predictions)

print("SVM Classification Report:")
print(svm_classification_report)
```

```
SVM Classification Report:
              precision    recall  f1-score   support

     3         0.00         0.00         0.00         5
     4         0.40         0.08         0.13        25
     5         0.62         0.57         0.60       291
     6         0.53         0.77         0.63       432
     7         0.65         0.25         0.36       192
     8         0.00         0.00         0.00        35

 accuracy          0.56        980
 macro avg         0.37         0.28         0.29        980
weighted avg         0.55         0.56         0.53        980
```

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

```
_warn_prf(average, modifier, msg_start, len(result))
```

```
In [2]: import pandas as pd
        from sklearn.model_selection import train_test_split
        from sklearn.preprocessing import LabelEncoder
```

```
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score

url = "https://raw.githubusercontent.com/tejavenkat473/Machine-Learning/main/winequality-white.csv"
data = pd.read_csv(url)

# Preprocessing
X = data.drop('quality', axis=1) # Assuming 'target_variable' is the column to predict
y = data['quality']

# Convert categorical variables to numerical using Label encoding
label_encoder = LabelEncoder()
X_encoded = X.apply(label_encoder.fit_transform)

# Create instances of the classification algorithms
rf = RandomForestClassifier()

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)

rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)
accuracy_rf = accuracy_score(y_test, y_pred_rf)
print("Random Forest Accuracy:", accuracy_rf)
```

Random Forest Accuracy: 0.6989795918367347

In [112...

```
# Calculate confusion matrix
rf_cm = confusion_matrix(y_test, y_pred_rf)

# Calculate F1 score
rf_f1 = f1_score(y_test, y_pred_rf, average='weighted')

print("Random Forest Confusion Matrix:")
print(rf_cm)
print("Random Forest F1 Score:", rf_f1)
```

Random Forest Confusion Matrix:

```
[[ 0  0  1  4  0  0]
 [ 0  5 13  7  0  0]
 [ 0  5 207 77  2  0]
 [ 0  0 62 345 25  0]
 [ 0  0  4  71 113  4]
 [ 0  0  1  9  9 16]]
```

Random Forest F1 Score: 0.6920988476095293

In [113...

```
# Generate classification report
rf_classification_report = classification_report(y_test, y_pred_rf)

print("Random Forest Classification Report:")
print(rf_classification_report)
```

Random Forest Classification Report:

	precision	recall	f1-score	support
3	0.00	0.00	0.00	5
4	0.50	0.20	0.29	25
5	0.72	0.71	0.72	291
6	0.67	0.80	0.73	432
7	0.76	0.59	0.66	192
8	0.80	0.46	0.58	35
accuracy			0.70	980
macro avg	0.57	0.46	0.50	980
weighted avg	0.70	0.70	0.69	980

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

In [114...

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder
```

```

from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score

url = "https://raw.githubusercontent.com/tejavenkat473/Machine-Learning/main/winequality-white.csv"
data = pd.read_csv(url)

# Preprocessing
X = data.drop('quality', axis=1) # Assuming 'target_variable' is the column to predict
y = data['quality']

# Convert categorical variables to numerical using Label encoding
label_encoder = LabelEncoder()
X_encoded = X.apply(label_encoder.fit_transform)

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X_encoded, y, test_size=0.2, random_state=42)

# Create an instance of the K-nearest Neighbors algorithm
knn = KNeighborsClassifier()

# Train the model
knn.fit(X_train, y_train)

# Make predictions on the test set
y_pred_knn = knn.predict(X_test)

# Evaluate the accuracy of the model
accuracy_knn = accuracy_score(y_test, y_pred_knn)
print("K-nearest Neighbors Accuracy:", accuracy_knn)

```

K-nearest Neighbors Accuracy: 0.5530612244897959

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\neighbors\\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

In [116...

```

# Calculate confusion matrix
knn_cm = confusion_matrix(y_test, y_pred_knn)

# Calculate F1 score
knn_f1 = f1_score(y_test, y_pred_knn, average='weighted')

print("KNN Confusion Matrix:")

```

```
print(knn_cm)
print("KNN F1 Score:", knn_f1)
```

KNN Confusion Matrix:

```
[[ 0  0  3  2  0  0]
 [ 0  4 14  6  1  0]
 [ 0  7 169 103 10  2]
 [ 0  4  97 280 47  4]
 [ 0  1  16  88 85  2]
 [ 0  0  2  15 14  4]]
```

KNN F1 Score: 0.5426100213715772

In [117...

```
# Generate classification report
knn_classification_report = classification_report(y_test, y_pred_knn)

print("knn Classification Report:")
print(knn_classification_report)
```

knn Classification Report:

	precision	recall	f1-score	support
3	0.00	0.00	0.00	5
4	0.25	0.16	0.20	25
5	0.56	0.58	0.57	291
6	0.57	0.65	0.60	432
7	0.54	0.44	0.49	192
8	0.33	0.11	0.17	35
accuracy			0.55	980
macro avg	0.38	0.32	0.34	980
weighted avg	0.54	0.55	0.54	980

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\metrics\\_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero\_division` parameter to control this behavior.

\_warn\_prf(average, modifier, msg\_start, len(result))



In [115...

```
# Create a list of classification models
models = [
    ('Logistic Regression', LogisticRegression()),
    ('Decision Tree', DecisionTreeClassifier()),
    ('Random Forest', RandomForestClassifier()),
    ('Support Vector Machine', SVC()),
    ('K-Nearest Neighbors', KNeighborsClassifier())
]

# Iterate over each model and print its accuracy
for name, model in models:
    model.fit(X_train_scaled, y_train)
    predictions = model.predict(X_test_scaled)
    accuracy = accuracy_score(y_test, predictions)
    print(f"{name} Accuracy:", accuracy)
```

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\linear\_model\\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
Logistic Regression Accuracy: 0.5306122448979592
Decision Tree Accuracy: 0.6173469387755102
Random Forest Accuracy: 0.7051020408163265
Support Vector Machine Accuracy: 0.5612244897959183
K-Nearest Neighbors Accuracy: 0.5428571428571428
```

C:\Users\Administrator\anaconda3\lib\site-packages\sklearn\neighbors\\_classification.py:228: FutureWarning: Unlike other reduction functions (e.g. `skew`, `kurtosis`), the default behavior of `mode` typically preserves the axis it acts along. In SciPy 1.11.0, this behavior will change: the default value of `keepdims` will become False, the `axis` over which the statistic is taken will be eliminated, and the value None will no longer be accepted. Set `keepdims` to True or False to avoid this warning.

```
mode, _ = stats.mode(_y[neigh_ind, k], axis=1)
```

In [ ]: